

DBMS

File Systems

- 1) data gets duplicated (data redundancy) which results in inconsistency.
 - 2) sorting of data is difficult
 - 3) Data retrieval is difficult
 - 4) more storage space
 - 5) more human intervention required
 - 6) sharing files is difficult, ensuring security is difficult
- Some data stored in many places.
if value of the same data stored in different places have different values

Different DBMS - Oracle
 - SQL Server } etc.
 - MS Access Excel

DBMS → collection of inter related data and a set of programs to access those data. (software used to create, manipulate and delete DB)
 - goal is to provide a way to store and retrieve DB info that is both convenient and efficient.

Major drawbacks of file processing systems

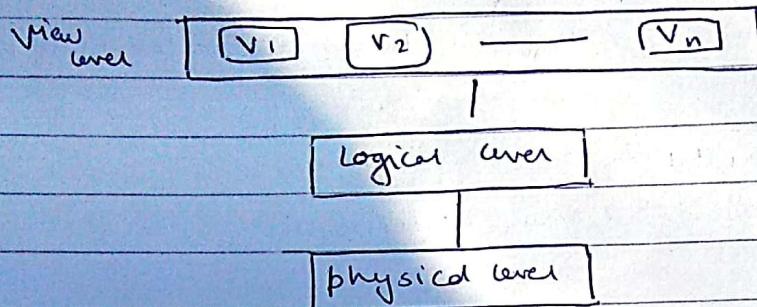
- 1) data redundancy and inconsistency
- 2) difficulty in accessing data
- 3) integrity problems (some constraints which should not be violated)
 - adding new constraints is difficult
- 4) Atomicity problem (eg: banking transaction)
 - all or none
 - ↳ internet
- 5) Concurrent Access Problem (eg: we are booking tickets at the same time in railways)
- 6) Security problem
 - ↳ who all can read, or who all can write.

View of Data

- 1) Physical level - how the data are actually stored
- 2) Logical level - what data are stored in DB and what relationship

exist among them.

- 3) View level - describes only part of the entire DB.
- simplify interaction with the system



Instances and Schema

- collection of info. stored in a DB at a particular moment - instance
- design of DB is called DB schema
instance can change, schema can not change
- DB systems have several schemes partitioned according to levels of abstraction.
 - 1) Physical Schema - physical interpretation and storage
 - 2) Logical Schema - DB design at logical level
 - 3) Subschemas - at the view level that describes

Data Models

- collection of conceptual tools for describing data, data relationships, data semantics and consistency constraints

- 1) Relational Model - uses collection of tables to represent both data and relationship among those data.

- each table has multiple columns and each column has a unique name
- table is also called relations

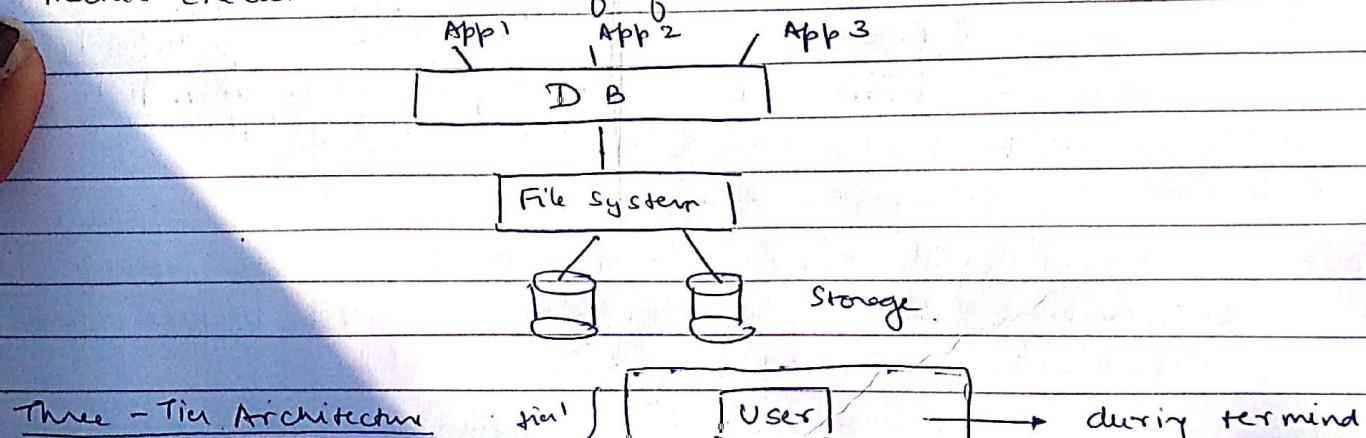
DBMS

2) Entity Relationship model (ER) -

- uses collection of basic objects called entities and relationships among these objects.
- an entity is a thing or object in the real world that is distinguishable from other objects.

Database Architecture

- influenced by the underlying computer system on which DB system runs.
- DB systems can be centralized or client-server where one server machine executes work on behalf of multiple client machines.



- 2) Designer chose a data model and translates these requirements into a conceptual schema of the DB.
The ER model is used to represent the conceptual design.
- 3) A fully developed conceptual schema indicates final requirements of the enterprise.
- 4) Process of moving from an abstract data model to the implementation of the DB proceeds in 2 final design phases.
 - a) in logical design phase - mapping conceptual schema into a relational schema
 - b) finally, designs uses the resulting system specific DB schema in the physical design phase.

Physical features of DB are specified, index, file organizations, etc.

2 major pitfalls

- 1) Redundancy -
- 2) Incompleteness - a bad design may make certain aspects of the enterprise difficult or impossible to model

Entity Relationship Model

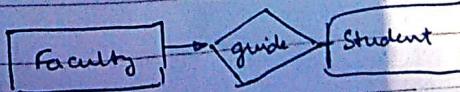
- 1) entity set - An entity is a thing or object in the real world that is distinguishable from all other objects.
- 2) An entity set is a set of entities of some type that share some properties or attributes
 - entity set do not need to be disjoint.
 - An entity is represented by a set of attributes e.g. id, name, dept, etc.

2) relationship sets -

- association among several entities

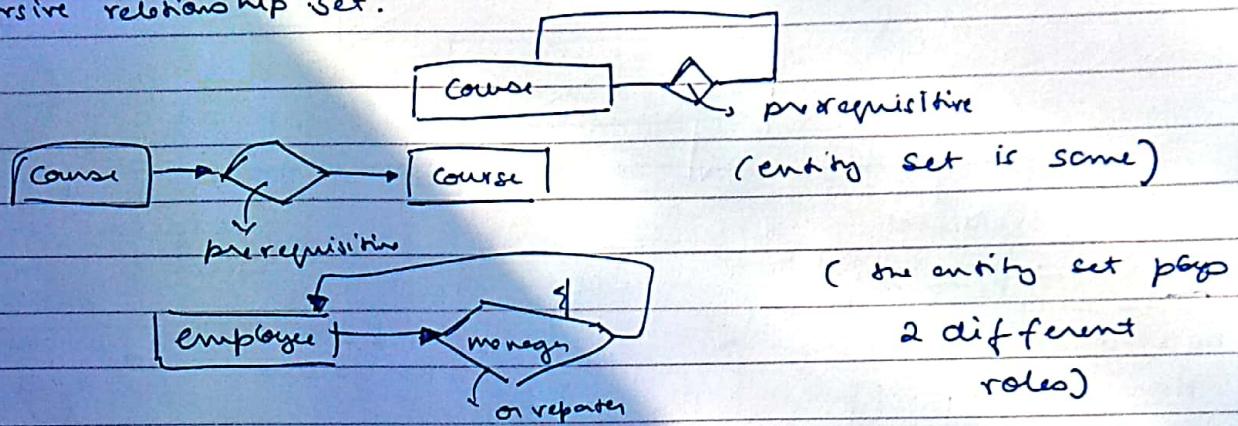


not necessary
to allow related
(some of them are selected)



DBMS

- Association between entity sets is referred to as participation
- The function that an entity plays in a relationship is called the entity role. (predefined)
- Recursive Relationship - some entity set participates in a relationship set more than once, in different roles called recursive relationship set.



- A relationship may also have attribute called descriptive attribute.

• Single valued and multi-valued

- only one value possible. eg: id, name, dob

single valued (normal representation) multi-valued } we need this for capturing info

age

• Derived Attribute eg: age can be calculated from date of birth.

represented as are not stored in the database at all.

• value of derived attribute is not stored at all.

• Null Values ⇒ Attributes can have NULL values if

- the entity does not have a value for it. eg: middle name
- value is missing/unknown

- a) missing - exist but does not have information
- b) unknown - do not know if information exist.

- Simple and Composite Attributes

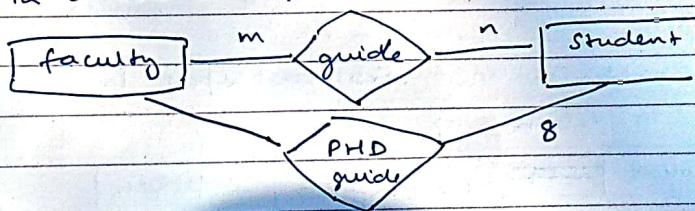
not divided into subparts → divide into subparts → eg: name (middle, first, last)
 eg: dob, pan no, etc. address

- Constraints - to which contents of a DB must conform

eg: in an email @ should be used. , phone no should be numeric and 10 digits
 number

1) Mapping Cardinalities (cardinality ratio)

- express no of entities to which another entity can be associated with via a relationship set.



Mapping Cardinalities

- useful in binary relationships
- in a binary relationship, mapping cardinalities must be one of the following:

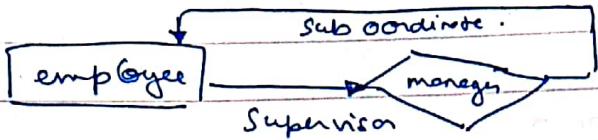
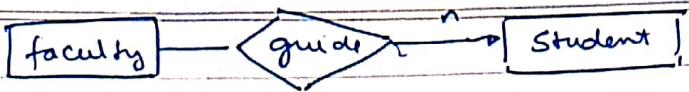
a) One-to-one



(there may be a new employee, which has not yet been assigned workstation or vice versa? → they do not participate in the relationship.)

b) One-to-many

- student can be guided by a single faculty
- faculty may guide multiple students.



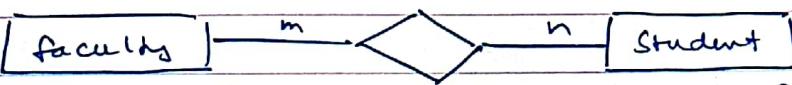
multiple employees can report to a same manager
(but they report to just one manager)

c) Many to one



a single state can have many cities, but a city can belong only to one state
~~(if we)~~ (backward representation of one to many)

d) Many to many



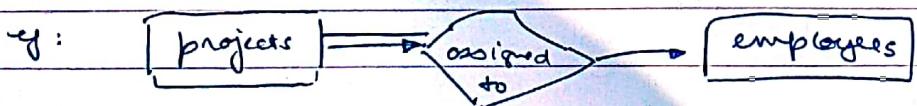
without the restriction that one student can be guided by only one faculty

Participation

if ~~one~~

(participation from the employees side is partial)

Total participation
if every entity in one entity set participate in at least one relationship in relationship set -



participation from the project side is total (symbolized by =)

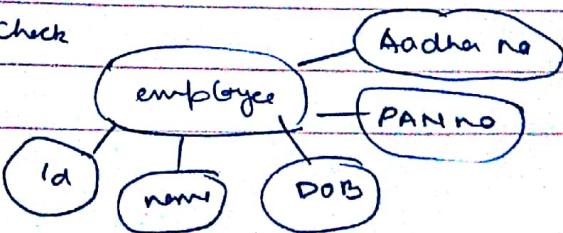
Keys - Keys for an entity is a set of attributes that suffices to distinguish entities from each other.

^{as it's compulsory}

x	x	x	unique & not null	can be null
id	name	DOB	DOJ	PAN.NO Aadhar NO
1	ABC		:	A 101
2		XYZ	:	B 102
3	ABC		:	C 103
4		XYZ		D 104

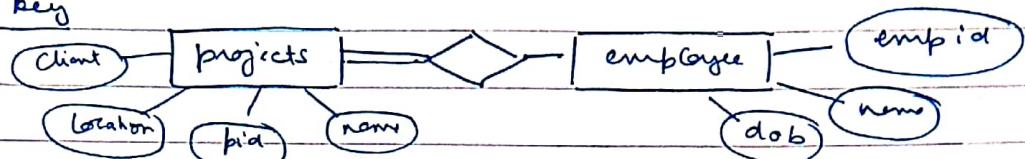
uniqueness and two things to check
not null

for keys



4 Types of Keys

- 1) Primary key \Rightarrow unique + not NULL represented by 
- 2) Candidate keys \Rightarrow which can be the primary keys
we have to choose one and make them primary key.
- 3) Alternate key \Rightarrow all the candidate keys except primary key
- 4) Foreign key



total participation from project side.

when we map ER model to table.

this is a primary key
of table projects

emp_id	name	- - -	bid

\Rightarrow it is a foreign key

foreign key can have null values and repeated values also.

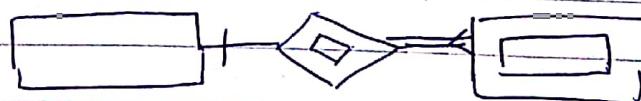
(it will have only values stored in the original table)

Weak Entity Set

- entity set that does not contain sufficient attribute to form a primary key is termed as a weak entity set.

Entries belonging to a weak entity set are identified by being related to specific entities from another entity set in combination with one of their attribute values.

even if a key contains all the components of the column, they will not be uniquely identified.



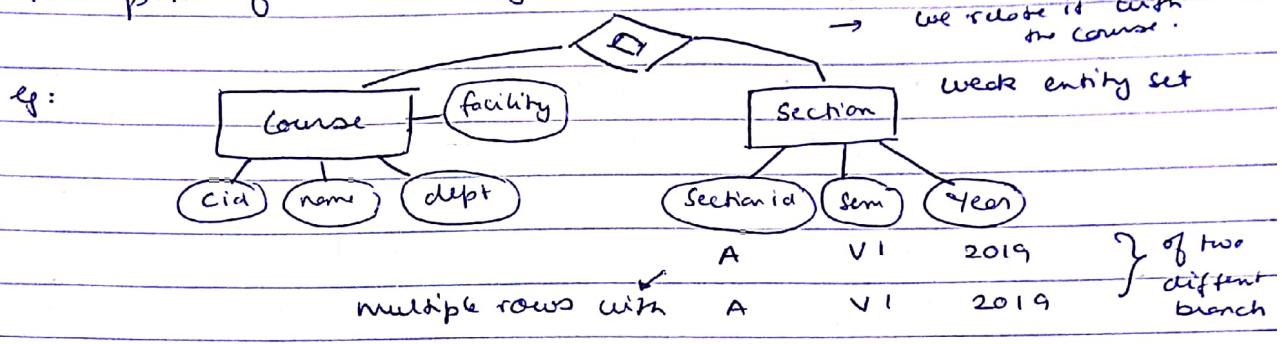
Strong entity set
(Identifying entity set)

Weak entity set

- This entity set is called the identifying or owner entity set.

DBMS

- Weak entity set must be associated with an identifying set.
- Identifying relationship set is many-to-one from weak entity set to the identifying entity set.
- Participation of the weak entity set in the relationship is total.



(cid will be used for identify uniquely)

of one set
data

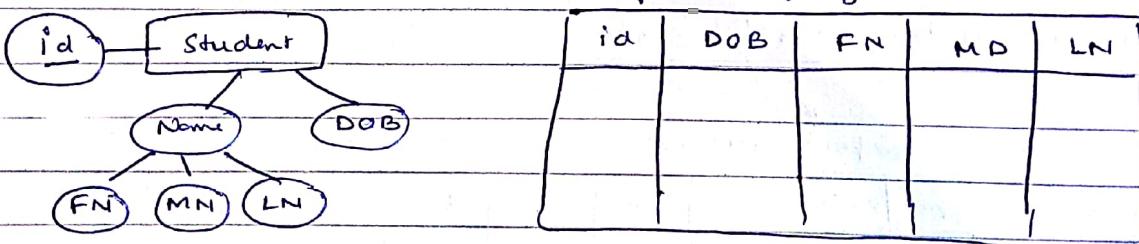
primary keys. (alternate keys) can not be used in another entity set.

Reduction to Relational Schema

- Representation of strong entity sets with simple attributes
- SES → table with attributes as columns

primary key of entity set → primary key of relationship / table.

+ primary key



for complex attributes

- For composite attributes → Create separate attributes for each of the component attributes (eg: name)
- derived attributes + not stored explicitly (eg: age)
- multivalued attribute. → (eg: phone no.)

two possibilities

S1	P1
S1	P2
S2	P3
S2	P4

Some data

but here data
redundancy is
there.

but we will make another table as.

for each multi value attribute.

Sid	Phone -
S1	P1
S1	P2
S2	P4
S2	P5

null's
will not
be here
If min no
of phone
no's is
3

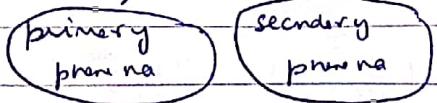
and if we have the minimum no of phone no.

for eg:

Sid	Phone 1	Phone 2	Phone 3

data redundancy will
not be there but here
the problem is NULL

this is not multivalued as each will only have
1 value.



Representation of WES

$$A \text{ (WES)} = \{a_1, a_2, \dots, a_m\}$$

attributes of the weak entity set.

$$B \text{ (IES)} = \{b_1, b_2, \dots, b_n\}$$

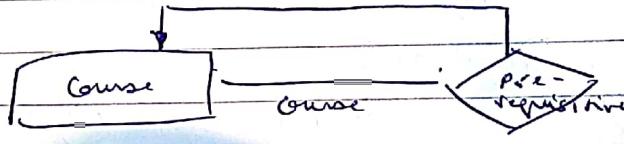
Composite primary key of the identifying entity set.

to	a ₁	a ₂	- - -	a _m	b ₁	b ₂	- - -	b _n
----	----------------	----------------	-------	----------------	----------------	----------------	-------	----------------

Representation of Relationship sets

$R = \{a_1, a_2, \dots, a_m\} \rightarrow$ union of primary keys of each of
the entity sets participating in R.

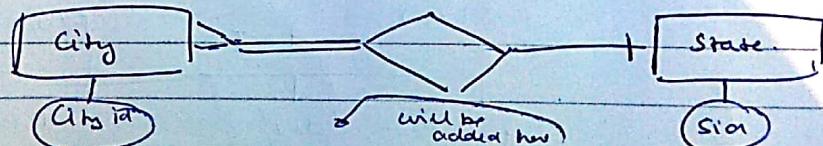
- $\{b_1, b_2, \dots, b_n\}$ - descriptive attribute.
- For Binary Relationship (primary key)
 - many-to-many
 - union of primary key of both the entity set.
 - one-to-one
 - primary key of either entity set can be chosen (as the primary key)
 - many-to-1 or 1-to-many relationships
 - primary key on the ES on the many side of the relationship set serves as the primary key.
- For n-ary relationships, union of primary key from participating Entity set is taken as the primary key of the table.
- For recursive relationship set, pre-required - roles indicators associated with the relationships are used as attribute name.
 (Combination of both will be used as primary key) PK combination of both



Course	Prereq

Identifying relationship \Rightarrow no scheme is created.

- Combination of schemes \Rightarrow
 - for many-to-1 relationship and participation $\#$ of A (many side) is total.
 - \Rightarrow combine A and AB to form a single scheme.



Cid	Name	---	Sid	Cid	Sid

(for many-to-one this should be)

and this is how it is stored

will be rejected.

D BMS

UNIT-2

Relational Algebra and Relational Calculus

- RA and RC were developed before SQL.
- The basic set of operations for the formal relational model is called the relational algebra.
- These operations enable a user to specify basic retrieval requests as relational algebra expressions.
- The result is a new relation.
- The relational calculus provides a higher level declarative language for specifying relational queries.
- In relational calculus, there is no order of operations to specify how to retrieve the query result - only what info the result should contain.

Relational algebra operations can be divided into 2 groups =>

- 1) include set operations - union, intersection, set difference, cartesian product
- 2) operations specifically for relational databases, e.g. SELECT, PROJECT, JOIN, etc.

Unary Relational Operations -

- 1) SELECT (σ) - to choose a subset of tuples (rows) from a relation that satisfies a selection condition.

$\sigma_{D.NO=5} (emp)$ emp = table or relation
↓
selection condition.

result in all the employees from the emp table whose department no is 5.

In SQL => $\text{SELECT } *$ * (all the column names)
 FROM EMP
 WHERE D.NO = 5

- Selection Condition is a boolean expression on the attributes of relation R (can be true or false)
- Clauses can be connected by standard boolean operators (and, or, not)

In case of Select, all rows will be distinct.

To form a general Selection Condition.

e.g: $\sigma_{(D.NO=5 \wedge Manager = 'Smith')} (emp)$

We want all employees whose dept no is 5 and manager is Smith.

- The Selection Condition is applied independently to each individual independently to each individual tuple in R.
- Select operator is unary - applied to a single relation
- degree (no of columns) of resulting from a select operation
= degree of the original relation
- The fraction of tuples selected by a selection condition is referred to as the selectivity of the condition.
- $\sigma_{<cond_1>} (\sigma_{<cond_2>} (R)) = \sigma_{<cond_2>} (\sigma_{<cond_1>} (R))$
 σ operator is commutative.

2) Project operation (Π)

- Selects certain columns from the table and discards other col's.
 - $\Pi_{<name, salary>} (emp)$
 - ↳ it will give all the rows, but only name and salary.
 - ↳ this will pull the entire data of employees who D.NO = 5
 - $\Pi_{<name, salary>} (\sigma_{(D.NO=5)} (emp))$
 - ↳ name and salary of those employees whose D.NO = 5
 - ↳ final outcome.
- Removes any duplicate tuples
- If duplicates are not eliminated, result would be a multi set
(DO NOT CONFUSE WITH SQL \Rightarrow it is different)
- Commutativity does not hold on PROJECT.

$\Pi_{name} (\Pi_{name, DOB} (emp))$
↳ It will result in name of all the employees.

In case of Select, all rows will be distinct

To form a general Selection Condition:

e.g.: $\sigma_{(D.NO=5 \wedge Manager = 'Smith')}$ (emp)

↳ we want all employees whose dept no is 5 and manager is Smith.

The selection condition is applied independently to each individual independently to each individual tuple in R.

Select operator is unary - applied to a single relation

- degree (no of columns) of resulting from a select operation
= degree of the original relation

- The fraction of tuples selected by a selection condition is referred to as the selectivity of the condition.

$$\sigma_{\text{Cond}_1} (\sigma_{\text{Cond}_2} (R)) = \sigma_{\text{Cond}_2} (\sigma_{\text{Cond}_1} (R))$$

σ operator is commutative.

2) Project operation (Π)

→ selects certain columns from the table and discards other coln.

- $\Pi_{\text{name}, \text{salary}} (\text{emp})$

↳ it will give all the rows, but only name and salary.

this will pull the entire data of employees who D.NO = 5

- $\Pi_{\text{name}, \text{salary}} (\sigma_{D.NO=5} (\text{emp}))$

↳ name and salary of those employees whose D.NO = 5

↳ find outcome.

- Removes any duplicate tuples

If duplicates are not eliminated, result would be a multi set

(DO NOT CONFUSE WITH SQL ⇒ it is different)

Commutativity does not hold on PROJECT.

$\Pi_{\text{name}} (\Pi_{\text{name}, \text{DOB}} (\text{emp}))$

↳ it will result in name of all the employees.

$\Pi_{\text{Name}, \text{emp}} \text{DOB} (\Pi_{\text{Name}}(\text{emp}))$ is not defined.

- In SQL, PROJECT attribute list is specified in the SELECT clause of the query.
- If we remove the keyword DISTINCT from SQL query, duplicates will not be eliminated.

SELECT Salary
FROM Emp

~ duplicates will not be removed

SELECT DISTINCT Salary
FROM Emp

if we want non-DISTINCT duplicate values, is used

Inline expressions

$\Pi_{\text{fname}, \text{name}, \text{sal}} (\sigma_{DNO = 5} (\text{emp}))$

another option

* but the output of this expression is a table and then apply the other operation on this table.

- Or explicitly show sequence of operations giving a name to each intermediate relations and using assignment operation (\leftarrow)

$D_{\text{emp}} \leftarrow \sigma_{DNO = 5} (\text{emp})$
 $\text{Result} \leftarrow \Pi_{\text{fname}, \text{name}, \text{sal}} (D_{\text{emp}})$

attribute names will be same as the original one
i.e. fname, name, sal.

If we want to change the attribute values do

$\text{Result} (\text{First name}, \text{Last name}, \text{Salary}) \leftarrow \Pi_{\text{fname}, \text{lname}, \text{sal}} (D_{\text{emp}})$

Rename operation (ρ)

$\rho_S (B_1, B_2, \dots, B_n) (R)$

$S \Rightarrow$ new relation name

$B_1, B_2, \dots, B_n \Rightarrow$ new attribute names γ in order

we can not explicitly change the values of attributes in between,
we will have to specify the ~~the~~ attribute names before it also.

- In SQL, use keyword AS

```
• SELECT E.Fname AS First-Name,  
      E.Lname AS Last-Name,  
      E.Sal AS Salary-of-emp  
  FROM Emp AS E  
 WHERE E.D.NO = 5
```

Relational Algebra Operations from Set Theory

- UNION, INTERSECTION, SET DIFF

- binary operations

- Union Compatibility \Rightarrow

2 relations must have same type of tuples

1) No of columns should be same

2) data type of each corresponding columns should be same.

EID	Ename	DOJ	PAN
1	John	2023-01-01	1234567890

Emp

S.No	Sname	DOB	PAN
1	Akash	2001-01-01	1234567890

Student

both these tables are union compatible.

(even though it is useless)

UNION (RUS)

- includes tuples that are either in R or in S or in both
- Duplicates are eliminated if the entire row is same
- Names of the column will be the names of the column

Union of Emp table and Student table.

EID	Ename	DOJ	PAN
1	John	2023-01-01	1234567890

? Data
in
emp

INTERSECTION Rns (Common ' between the two tables)

SET-DIFF (R-S) all the rows that are in R but not in S

(Emp - Student) will result in all the employees which are exclusively employees (i.e. they are not students)

CARTESIAN PRODUCT (no concept of union compatibility)

(Cross product) Cross Join X

Emp x Student =>

ENO	ENAME	DOS	PAN	SNO	STNAME	DOB	PAN
E1					S1		
E1					S2		
E1					S3		
E2					S1		
E3							

R (A₁, A₂ — A_n) × S (B₁, B₂ — B_m)

= Φ (A₁, A₂ — A_n, B₁ — B_m)

no of columns of both the table should not be same
may

Join (X) combines related tuples from 2 relations into single layer tuple.

Join ≈ Cartesian Product (Selection condition)

Emp |x| Dept

Gmp. D.NO = Dept D.NO

First it will Create a Cartesian

product and then choose
to all those values where d.no
= dept d.no.

R |x| <join condition> S

Result of join is a relation e.g. Q with n+m attributes

INTERSECTION Rns (Common between the two tables)

SET-DIFF (R-S) all the rows that are in R but not in S

"(Emp - Student) will result in all the employees which are exclusively employees (i.e. they are not students)

CARTESIAN PRODUCT (no concept of union compatibility)
(Cross product) Cross Join X

Emp \times Student =?

ENO	ENAME	DOS	PAN	SNO	STNAME	DOB	PAN
E1				S1			
E1				S2			
E1				S3			
E2				S1			
E3							

R (A₁, A₂ — A_n) \times S (B₁, B₂ — B_m)

= Φ (A₁, A₂ — A_n, B₁ — B_m)

no of columns of both the table should not be same

Join (Ix) Combines related tuples from 2 relations into single larger tuple.

Join \approx Cartesian product (Selection condition)

Emp Ix Dept

Emp.D.NO = Dept.D.NO

First it will create a Cartesian product and then choose all those values where d.no = dept.d.no.

R Ix <join condition> S

Result of join is a relation P, Q with n+m attributes

$R(n)$

$S(m)$

- Q has 1 tuple for each combination of tuples
- A join with such a general condition is called theta join.
- Tuples whose join attributes are NULL or for which the join condition is false do not appear in the result.

Equi join or Natural Join

- "only = is used"
- have one or more pairs of attributes that have identical values in every tuple.
∴ one of each pair of attribute with identical value is superfluous
- Natural Join = equi join + removal of superfluous attribute
- Natural Join requires that the two join attributes have the same name in both relations.
- Join Selectivity - rate of expected size of join result divided by the minimum size $n_R \times n_S$

Inner join - match and combine operation defined as a combination of Cartesian product and selection.

$$- R \bowtie S = \sigma_{\text{condition}}(R \times S)$$

Division Operation (\div)

- Result of division is a relation $T(Y)$ that includes a tuple t if tuples t_R appears in R with $t_R[Y] = t$ and $t_R[X] \in t$

Retrieves names of employees who work on all the projects that 'John Smith' works on.

$$\text{SMITH} \leftarrow \sigma_{Fname = 'John', LName = 'Smith'}(Emp) \rightarrow \begin{array}{l} \text{all dept} \\ \text{employees} \\ \text{named Smith} \end{array}$$

$$\text{SMITH} - PNO \bowtie T_{PNO}(\text{WORKS-ON} \bowtie \text{SMITH})$$

It will create join of tables work-on and SMITH.
Grp.NO = PNO.

→ this picks all the project nos when John Smith is working on.
works on table

Emp table					
G.NO	E.Name	DOB	DOJ	PAN	
E1					
E11	John Smith				
E2					

this data will be produced by

query 1 resulting in
table Smith

G.no	P.no	
E1	P1	
E1	P6	
E11		
E11	P24	
E11	P12	
E11	P6	

E11	John Smith	.	-	-
		SMITH		

these two would now be joined
when E.WO = Emp NO

Output no-w will be.

PNO
P24
P12
P6

SMITH - PNOs

b) Emp-Smith (ENO) ← Works-On ÷ Smith-PNOs

↓

it will pull all the data of all employees (one by one)
and it will divide the value by the result table Smith-PNOs
(if exists completely or not)

Result ← $\prod_{\text{from}} \text{ENO}, \text{Name} (\text{Emp-Smith} \times \text{Emp})$

- SQL - exists

True
False

1) Generalized Projections

extends projection operation by allowing functions of attributes to be included in the projection list.

$\prod_{\text{ENO}}, \text{Sal} (\text{Employee})$

$\prod_{\text{ENO}}, \text{Sal} - \text{OIST} + \text{Sal} (\text{EMP})$

Operations can take place in generalized projections

DBMS

2) Aggregate functions

- Common functions include SUM, AVERAGE, MAXIMUM, MINIMUM
for SUM to function properly, data type must be integer.

display all those departments whose average salary is greater than 50,000

.Emp (eno, ename, Dno, sal)

Dept (dno, dname, loc)

Q/P \Rightarrow Average (dno, dname) > 50,000

SELECT TT_AVERAGE (sal) (Emp) } it will compute the average of all the employees of every department

TT_dno (dept no) } contains all the department no

3) GROUPING

- grouping tuples in a relation by the value of some of their attributes and then applying an aggregate function independently to each group.

P \downarrow
 $\text{PR}(\text{DNO, no of emp, avg sal})$ (D.NO) } count (emp) & average (sal)
 this will give the answer for the above question.
 it will result in a table.

D.NO.	no-of-employees	av.sal
-------	-----------------	--------

- In general, duplicates are not eliminated when an aggregate function is applied.
- NULL values are not considered in the aggregation.

Recursive Closure Operations

- applied to recursive relationship between tuples of some type.
e.g.: relationship between employee and manager

- Report all reportees of an employee e^* at all levels
i.e. all emp e^1 directly supervised by e^*
and all emp e^{11} directly supervised by e^1 and so on.

Eid	Ename	Bno	Sid
E1			E10
E10			E34
E34			E107
E107			X
E2			E34
E202			E34

direct
 E_{34} direct reportees
 E_{10}, E_2, E_{202}

and then it will pull all the direct reportees of E_{10}, E_2, E_{202}

- easy to specify all employees supervised by e at a specific level by joining table with itself one or more times.

However, difficult to specify all supervisors at all levels.

$\text{James_id} \leftarrow \pi_{\text{eid}} (\sigma_{\text{name} = \text{James}} (\text{emp}))$

$\text{Supervision}(\text{eid}_1, \text{eid}_2) \leftarrow \pi_{\text{eid}_1, \text{sid}} (\text{emp})$

$\text{Result}(\text{eid}) \leftarrow (\text{Supervision} \bowtie_{\text{sid} = \text{eid}_2} (\text{James}))$

Supervision	
Eid1	Eid2
E1	E10
E2	E34
E10	E34
E34	E107
E107	-
E202	E34

James_id

eid
E34

For second level

$\text{RESULT2}(\text{eid}) \leftarrow \pi_{\text{eid}_1} (\text{Supervision} \bowtie_{\text{eid} = \text{eid}_2} \text{Result})$

Outer Join Operations

- Join Operation match tuples that satisfy the join condition
- Type of join where tuples (rows) without match are eliminated

is known as inner join.

Outer join → user wants to keep all the tuples in R or all those in S, or all those in both relations in the result of the JOIN regardless of whether or not they have matching tuples in other relations.

→ $R \bowtie S$ left Outer join ⇒ all the rows of R must be included irrespective of whether they satisfy the condition or not.

→ $R \bowtie S$ right Outer join (all the rows of S are included)

→ $R \bowtie S$ full outer join

- list of all emp names as well as names of the department they manage if they happen to manage a dept. If they do not manage one, we can indicate with a NULL value.

Emp			Dept		
Eid	ename	-	Dno	ename	Mg_id
E1	-		D1		E2
E2	-		D2		E3
E3	-		D106		E2

Expl / DS / Dept \bowtie / first part

T ename, deptname (Emp \bowtie Dept)

eid = mgrid

→ for all the managers of the department.

T ename, deptname (Emp \bowtie Dept)

eid = mgrid

Outer Union operation

→ to take union of tuples from 2 relations that have some common attribute but are not union compatible.

R (x, y) and S (x, z)

Y columns of table R

R and S are tables

• Tuples are partially compatible

- attributes that are union compatible are represented only once in the result, and those attributes that are not union compatible from either relation are also kept in the result relation.

$T(x, y, z)$

SQL (Structured Query Language)
 non procedural language (declarative language)

- SQL is classified as
 - 1) Data Definition Language (DDL) - Create, Alter, Drop, truncate
 - 2) Data Modification Language (DML) - Insert, Update, delete
 - 3) Data Control Language (DCL)
 - 4) Transaction Control Language (TCL)

DDL Operations are Auto Commit (Operations can not be rolled back)
 table once created can be dropped but the CREATE TABLE Command
 can not be rolled back.

DROP table deletes the entire table, whereas in TRUNCATE deletes the
 entire contents of the table, the table still remains.

Select is not ~~not~~ a DDL or DML ? it is just used for query

DML statements are not Auto Commit.

even delete operations can be undone. (it can be rolled back)

DB MS

- process of efficiently organizing data in a DB.
- a 2 step process =>
 - puts data into tabular form by removing repeating groups
 - removes duplicated data from the relational table.

Anomalies =>

- insertion anomaly
- updation anomaly
- deletion anomaly

Insertion anomaly (leads to updation anomaly)

emp

Gid	Fname	Lname	Pid	Pname	Ploc	→ split it into two tables
Ex	-	-	NULL	NULL	NULL	

Suppose an employee X has been assigned recruited but has not yet been assigned some project

the pid, pname, ploc will all be NULL in this case

1 or 2 columns are acceptable, but here we have more columns with NULL value.

Updation anomaly

→ happens when data is redundant.

(values are repeated many times)

Gid	Fname	Lname	City	Pincode
-			Bombay	
-			Bombay	



if name of city changes from Bombay to Mumbai

Some one has changed the city from Bombay to Mumbai

Others have not

→ Updation anomaly

Deletion anomaly

emp	fname	lname	Did	Dname
G12			D2	

→ if there is only one employee of a department and we delete that, the information pertaining to the departments is also lost forever.

Normal forms

A table is said to be in a normal form if it satisfies a set of constraints.

6 Normal forms

- 1) First Normal Form (1NF)
- 2) Second " " (2NF)
- 3) Third " " (3NF)
- 4) BCNF
- 5) Fourth (4NF)
- 6) ~~5th~~ 5 NF

item no	item-desc	price	manufacturer name	distributor id	distributor name	order id	date entered	qty
	Watches		-					

There may be multiple orders of watches from the same distributor

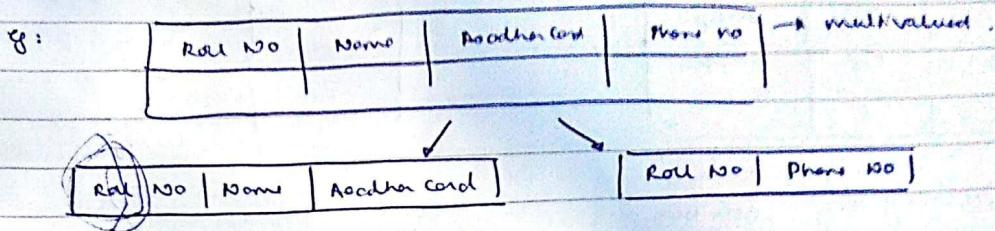
First Normal Form

No repeating groups → composite values
and
other all attributes are dependent on PK

multiple values

composite values

When we have a multivalued attribute we create another table.



all attributes must be dependent on primary key. (by default satisfied)

if we write $\begin{array}{l} \text{SELECT } * \\ \text{FROM PROJECT} \\ \text{WHERE Pid = 'P01'} \end{array}$] we will get a single value.

if value of project id (primary key) is known, all other values can be found out, and it will be unique. \rightarrow functionally dependent

Second Normal form (2NF)

if a table is in INF and all attributes are fully functionally dependent on all the candidate keys.

Fully functionally dependent

for a given relation R, attribute B of R is fully functionally dependent on attribute A of R, if it is functionally dependent on A and not on any subset of A.

Order - details

Oid	Prid	Prname	Manufacturer date	Price
O1	P1	Soap		
O1	P10	Shampoo		
O1	P6	Bread.		

Mfg date
and price
are fully
functionally
dependent

primary key
non key
attribute.

Prname is not dependent on Oid and Prid.

it ~~only~~ only needs Prid. to determine its value uniquely

it is not fully dependent. (it is dependent only on
functionally subset of primary key).

is not in

We need to split this table, because, this table is in 2NF.

(data is redundant over here)

so bring it

now 2NF of this table is.

Oid	Prid	Mfg date	Price	and	Prid	Prname
↑ primary key		(now primary key is of both is not composite)			↑ primary key	

now they are in 2 NF

Transaction

C#	Cname	C City	C phone	P #	Date	Quantity
----	-------	--------	---------	-----	------	----------

C \Rightarrow customer P \Rightarrow product.

Cname, C City, C phone depends on C#

C# + Date and time \Rightarrow primary key \Rightarrow

we will get all other columns (name, city, phone,

P#

f

2nd normal form of the table

C#	Cname	Cphone	Ccity
----	-------	--------	-------

C#	P#	Part name	Quantity
----	----	-----------	----------

Transitive Dependency → column

- value in a non-key field is determined by value in another non-key field and that field is not a candidate key.

$$(a,b) (b,c) \Rightarrow (a,c)$$

non-key

non-key

primary key →

$$A \rightarrow B$$

$$A \rightarrow B$$

$$PK \rightarrow A$$

$$\Rightarrow PK \rightarrow B$$

↑ primary key.

Project No	Project Title	Project Manager	Phone
------------	---------------	-----------------	-------

a project can
have only
one manager.

Project Mgr → Phone

Project No → Project Title

as Project No is

Project No → Project Mgr

a primary key.

Project No → Phone

there is transitive dependency.

table is in 2 NF and

all its attributes are dependent on ~~key~~ only on the

key attributes (primary key / candidate key)

Project Mgr is a non key attribute. ⇒ there is transitive

⇒ the table is not in 3NF. dependency.

key attributes should be mutually independent.

in table.

Project No	Project Title	Project Mgr
------------	---------------	-------------

Project Mgr	Phone
-------------	-------

2nd normal form of the table

C#	Cname	Cphone	Ccity
----	-------	--------	-------

CH	PT	Dot and time	Quantity
----	----	--------------	----------

Transitive Dependency, column

value in a non-key field is determined by value in another non-key field and that field is not a candidate key.

$$(a,b) \quad (b,c) \Rightarrow (a,c)$$

$$\begin{array}{l} \text{non-key} \\ \downarrow \\ A \rightarrow B \end{array} \quad \begin{array}{l} \text{non-key} \\ \downarrow \\ B \rightarrow C \end{array} \quad \begin{array}{l} \text{primary key} \Rightarrow \\ A \rightarrow C \\ PK \rightarrow A \\ \Rightarrow PK \rightarrow B \end{array}$$

↑ primary key.

Project No	Project Title	Project Manager	Phone
------------	---------------	-----------------	-------

a project can have only one manager.

Project Mgr → Phone

Project No → Project Title

as Project No is

Project No → Project Mgr

a primary key.

Project No → Phone

there is transitive dependency.

Table is in 2NF and

all its attributes are dependent on ~~key~~ only on the key attributes (primary key / candidate key)

Project Mgr is a non-key attribute. so there is transitive

⇒ the table is not in 3NF. dependency.

Non-key attributes should be mutually independent.

In the table,

Project No	Project Title	Project Mgr
------------	---------------	-------------

Project Mgr	Phone
-------------	-------

Normalization questions

(1) a) we have composite branch address

which can be broken ~~into~~ into

B.NO	B. A Lin 1	B. A Lin 2	B. Add L 3	City	State	Pincode
						↓ primary key

Telephone is a multivalued column, we create one more table.

B.NO.	B.Tel	→ primary key
-------	-------	---------------

it is in first normal form

Since primary key is not composite.

⇒ no case of functional dependency

it is in 2nd normal form

now for 3rd normal form ⇒ check for transitivity ⇒ dependency

$$A \rightarrow B, (B) \rightarrow C \Rightarrow A \rightarrow C$$

↓ if B is a non-key attribute.

City Pincode → City

Pincode → State

now to bring it into 3rd normal form table is now ~~given~~ ^{is}

↓ primary key

↓ foreign key

B.NO.	B AL1	B AL2	B AL3	Pincode	City State
					↓ primary key

Pincode	City	State
---------	------	-------

B.NO	B.Tel	→ primary key
------	-------	---------------

↓ foreign key

(2) Primary key Staff no # branch no ⇒ Name / week

↓ primary key

B.NO → B Add

Staff no → Name, Position

↓

(when position of staff does not change irrespective of branch)

⇒ functional dependency on primary key

part No

not in 2nd normal form

↑ pk	↑ pk
B.NO B.Add	Staff No Name Home Position
Staff No B.NO hours/week	

transitive dependency \Rightarrow all 3 are =
 or it is not in 3rd normal form.

Now: one staff is working in different branches at different position.

B NO \rightarrow B Add

staff NO \rightarrow Name

staff NO, B NO \rightarrow hours / week, position

B NO B Add

(3) Branches \rightarrow B Add, Tel NO, Mgr Staff No, Name.
 Manager can supervise many branches

B NO is the primary key.

B NO \rightarrow B Add, Tel NO, Mgr Staff No, Name.

Mgr Staff NO \rightarrow Name

\Rightarrow it is not in 3rd normal form

B.NO B.Add Tel NO SMS NO.

SMS NO Name

DBMS

18/02

Integrity Constraints

Create table emp

(name varchar not null ,

PAN , varchar unique

active . varchar default 'Y' ,

check in ('Y', 'N')

)

all the values with default should be towards the right end, otherwise

there will be confusion as to which the value has to be mapped.

otherwise we have to do explicit mapping.

insert into emp values ('ABC', 'A X L 1231P')

name

PAN

and active by

default 'Y' .

Check can be used to check whether the tuple values values of the columns should be from the specified one. (example for active check the column can contain only 'Y' and 'N' .

Referential Integrity

(for foreign key)

D.NO.	D.Name	Loc.		E.NO	E.Name	D.NO (F.K.)

Dept

Emp

If Department nos is D₁ to D₁₀ in Dept table

=> Dept No of Emp tabe should also be between D₁ to D₁₀

Any other dept is not valid.

Dept nos should be from the Dept table only. (but we can have NULL in the Emp table)

in the D.NO of Emp table)

If there are many Dept values of D₂ in the Emp table.

and if D₂ is deleted from Dept table

all the values with D₂ department in Emp table will not be valid anymore.

1 option is to assign department value to be NULL or

(in Emp table)

other option is to delete all the values with department D₂.

DBMS

18/02

Integrity Constraints

Create table emp

```
( name varchar not null,
  PAN , varchar unique
  active . varchar default 'Y',
  )
```

• primary key • not null • unique

• default

• check

check in ('Y', 'N')

all the values with default should be towards the right end, otherwise there will be confusion as to which the value has to be mapped.
otherwise we have to do explicit mapping.

insert into emp values ('ABC', 'A1231P')

name → PAN and active by
default 'Y'.

Check can be used to check whether the ~~tuple values~~ values of the columns should be from the specified one. For example for active check the column can contain only 'Y' and 'N'.

Referential Integrity

(for foreign key)

D. NO.	D. Name	Loc.		E. NO.	E. Name	D. NO (F.K.)

Dept

Emp

If Department nos is D₁ to D₁₀ in Dept table

⇒ Dept No of Emp table should also be between D₁ to D₁₀

Any other dept is not valid.

Dept nos should be from the Dept table only. (but we can have NULL in the Emp table)

in the D. NO of Emp table)

If there are more values of D₂ in the Emp table.

from Dept table

in the Emp table will not be valid anymore.

to be NULL or (in Emp table)
values with

Triggers

(action leads to immediate reaction)

If D6 is changed to D24 ~~then in the Dept table.~~

then the same should be changed for all rows with D6 department
in Emp table also

- Create Trigger => used to implement actions in SQL

create trigger <Trigger-name>

Before insert or update of salary, supervisor-id, on Emp
on after update salary = 0 where supervisor-id = 134

Trigger is regarded as event - condition - action (ECA) has 3 components.

1) Event(s) - usually DB update operations that are
explicitly applied to the DB *(insertion, deletion or update)*

~~→~~ - specified after keyword Before / After

2) Condition - determines whether the rule action should be executed

Once the triggering event has occurred, an optional condition may be evaluated. If no condition is specified, action will be executed once the event occurs.

- If a condition is specified it is evaluated first, and only if it evaluates to true will the rule action be executed.

The condition is specified in the where clause of the trigger

3) Action to be taken, usually a sequence of SQL statements.

Tuple Relational Calculus

Query Lang

(Relational algebra) Procedural

Non-Procedural

Relational algebra

Relational calculus

SQL

Relational Calculus

Tuple Relational
calculus

Domain Relational
calculus

- user is concerned with the details of how to obtain the end results
- uses mathematical predicate calculus
- a predicate is a truth valued function with arguments.
- when we replace with values for the arguments, the function yields an expression called a proposition, which will be either true / false.

Tuple Relational Calculus

- write one declarative expression to specify a retrieval request.
- The expressive power of relational algebra and relational calculus is identical

Tuple variables and Range relations

- TRC is based on specifying a no of tuple variables
- each tuple variable usually ranges over a particular DB relation.
- the variable may take as its value any intended individual tuple from that relation.

→ $\{t \mid \text{cond}(t)\}$ → t variable should satisfy the condition.

$\{t \mid \text{Emp}(t)\}$ ← ~~for all the~~ all those rows which belong to Emp table.

$\{t \mid \text{Emp}(t) \text{ AND } t.\text{sd} > 50000\}$

all those employees (entire row) of the employees whose salary is greater than 50000

$\Rightarrow t.\text{sal} \approx t[\text{sal}]$

You need to specify the following information in a TRC expression ⇒

1) to each tuple variable t , the range relation R of to $R(t)$ because if we specify only $\{t \mid t.\text{sd} > 50000\}$, it will scan all the info ~~unnecessary~~ of the entire database.

2) a condition to select particular combination of tuples.

3) Requested attributes. (we could retrieve only the specified columns)